

October 3, 2018

Remarks

I've put a lot of time into this project, sadly, to no avail. My best results are from one of my early adaptations of the MNIST assignment where my code is incredibly ugly and I'm fairly sure batch normalization is implemented incorrectly.

One of the limiting factors in my explorations this week has honestly been my computing power. I just don't have time to run training with an insane number of parameters on my laptop. Starting next week I'll (somehow) be running my code on something more capable.

Results

Cifar10:

- Top-1 Accuracy: 70%

Cifar100:

- Top-1 Accuracy: 18%
- Top-5 Accuracy: 40%

Program: cifar10v2.py

```
1 from fcns import *
2 import numpy as np
3 from tqdm import tqdm
4 import tensorflow as tf
5 import matplotlib.pyplot as plt
6
7 logs_path = "./tf_logs/"
8
9 # hyper-parameters
10 input_size = 32 # 32x32 imgs
11 num_channels = 3 # RGB
12 num_classes = 10
13 dim_layer = [num_channels, 32, 32, 32, 32, 32, 32, 64, 64, num_classes]
14 cp = { # conv parameters
15     1:{
16         'k':5, # conv window size (kxk)
17         'p':2, # pool window size (pxp)
18         'ks':2, # conv stride
19         'ps':1, # pool stride,
20         'kpad':"SAME",
21         'ppad':"SAME"
22     },
23     2:{
24         'k':3, # conv window size (kxk)
25         'p':2, # pool window size (pxp)
26         'ks':1, # conv stride
27         'ps':2, # pool stride,
28         'kpad':"SAME",
29         'ppad':"SAME"
30     },
31     3:{
32         'k':3, # conv window size (kxk)
33         'p':2, # pool window size (pxp)
34         'ks':1, # conv stride
35         'ps':1, # pool stride,
36         'kpad':"VALID",
37         'ppad':"VALID"
38     },
39     4:{
40         'k':3, # conv window size (kxk)
41         'p':2, # pool window size (pxp)
42         'ks':1, # conv stride
43         'ps':1, # pool stride,
44         'kpad':"VALID",
45         'ppad':"VALID"
46     },
47     5:{
48         'k':3, # conv window size (kxk)
49         'p':2, # pool window size (pxp)
50         'ks':1, # conv stride
51         'ps':1, # pool stride,
52         'kpad':"SAME",
53         'ppad':"SAME"
54     },
55     6:{
56         'k':3, # conv window size (kxk)
57         'p':2, # pool window size (pxp)
58         'ks':1, # conv stride
59         'ps':1, # pool stride,
60         'kpad':"SAME",
61         'ppad':"SAME"
62     }
63 }
64
65 d1 = dim(input_size, cp[1])
66 d2 = dim(d1, cp[2])
67 d3 = dim(d2, cp[3])
68 d4 = dim(d3, cp[4])
69 d5 = dim(d4, cp[5])
70 d6 = dim(d4, cp[6])
71 print(d1, d2, d3, d4, d5, d6)
```

```

72
73 BATCH_SIZE = 300
74 NUM_EPOCHS = 30
75 learning_rate = .1
76 display_epoch = 1
77
78 # regularization parameters
79 l2_lambda = .01/(sum([dim**2 for dim in dim_layer]))
80 print(l2_lambda)
81 drop_out = .9
82
83 def main():
84     class Data(object):
85         def __init__(self):
86             np.random.seed(31415)
87             ([self.x_train, self.x_val, self.x_test],
88              [self.y_train, self.y_val, self.y_test]) = \
89                 loadData("./cifar10_data")
90
91             self.index = np.arange(self.x_train.shape[0])
92
93         def get_batch(self):
94             choices = np.random.choice(self.index, size=BATCH_SIZE)
95             return self.x_train[choices,:,:,], self.y_train[choices,:]
96
97     x = tf.placeholder(tf.float32, [None,input_size,input_size,num_channels])
98     y = tf.placeholder(tf.float32, [None,num_classes])
99     keep_prob = tf.placeholder(tf.float32) # for dropout
100
101     # f:R28x28 -> R10
102     def f(x):
103         layer_1 = tf.nn.dropout( conv_layer(x, W[1], b[1], cp[1]), keep_prob )
104         layer_2 = tf.nn.dropout( conv_layer(layer_1, W[2], b[2], cp[2]), keep_prob )
105         layer_3 = tf.nn.dropout( conv_layer(layer_2, W[3], b[3], cp[3]), keep_prob )
106         layer_4 = tf.nn.dropout( conv_layer(layer_3, W[4], b[4], cp[4]), keep_prob )
107         layer_5 = tf.nn.dropout( conv_layer(layer_4, W[5], b[5], cp[5]), keep_prob )
108         layer_6 = tf.nn.dropout( conv_layer(layer_5, W[6], b[6], cp[6]), keep_prob )
109         layer_7 = tf.nn.dropout( fc_layer(layer_6, W[7], b[7]), keep_prob )
110         layer_8 = tf.nn.dropout( fc_layer(layer_7, W[8], b[8]), keep_prob )
111         return tf.squeeze(
112             tf.add( tf.matmul(layer_8, W['out']), b['out'] )
113         )
114
115     # Store layers weight & bias
116     # default dtype=float32
117     # WEIGHTS
118     W = {
119         1: tf.Variable(tf.random_normal( [cp[1]['k'],cp[1]['k'], num_channels, dim_layer
120             [1]] )),
121         2: tf.Variable(tf.random_normal( [cp[2]['k'],cp[2]['k'], dim_layer[1], dim_layer
122             [2]] )),
123         3: tf.Variable(tf.random_normal( [cp[3]['k'],cp[3]['k'], dim_layer[2], dim_layer
124             [3]] )),
125         4: tf.Variable(tf.random_normal( [cp[4]['k'],cp[4]['k'], dim_layer[3], dim_layer
126             [4]] )),
127         5: tf.Variable(tf.random_normal( [cp[5]['k'],cp[5]['k'], dim_layer[4], dim_layer
128             [5]] )),
129         6: tf.Variable(tf.random_normal( [cp[6]['k'],cp[6]['k'], dim_layer[5], dim_layer
130             [6]] )),
131         7: tf.Variable(tf.random_normal( [d6*d6*dim_layer[6], dim_layer[7]] )),
132         8: tf.Variable(tf.random_normal( [dim_layer[7], dim_layer[8]] )),
133         'out': tf.Variable(tf.random_normal( [dim_layer[8], num_classes] ))
134     }
135
136     # BIASES
137     b = {
138         1: tf.Variable(tf.random_normal([dim_layer[1]])),
139         2: tf.Variable(tf.random_normal([dim_layer[2]])),
140         3: tf.Variable(tf.random_normal([dim_layer[3]])),
141         4: tf.Variable(tf.random_normal([dim_layer[4]])),
142         5: tf.Variable(tf.random_normal([dim_layer[5]])),
143         6: tf.Variable(tf.random_normal([dim_layer[6]])),
144         7: tf.Variable(tf.random_normal([dim_layer[7]])),
145         8: tf.Variable(tf.random_normal([dim_layer[8]])),

```

```

139         'out': tf.Variable(tf.random_normal([num_classes]))
140     }
141
142     # models
143     logits = f(x)
144     prediction = tf.nn.softmax(logits)
145     correct_pred = tf.equal(tf.argmax(prediction, 1), tf.argmax(y, 1))
146     accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32))
147
148     # binar cross entropy loss with L2 penalty on weights
149     loss = tf.reduce_mean( tf.losses.softmax_cross_entropy(y, logits) ) + \
150         l2_lambda*tf.reduce_sum(
151             [tf.nn.l2_loss(var) for var in
152              tf.get_collection(tf.GraphKeys.TRAINABLE_VARIABLES)]
153         )
154     optim = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(loss)
155     init = tf.global_variables_initializer()
156
157     # Create a summary to monitor cost tensor
158     tf.summary.scalar("loss", loss)
159     # Create a summary to monitor accuracy tensor
160     tf.summary.scalar("accuracy", accuracy)
161     # Merge all summaries into a single op
162     merged_summary_op = tf.summary.merge_all()
163
164     with tf.Session() as sess:
165         data = Data()
166         sess.run(init)
167         # op to write logs to Tensorboard
168         summary_writer = tf.summary.FileWriter(logs_path,
169             graph=tf.get_default_graph())
170         # training
171         for epoch in range(NUM_EPOCHS):
172             avg_cost = 0.
173             num_batches = int( np.ceil( data.index[-1] / BATCH_SIZE ) )
174
175             for i in tqdm(range(num_batches)):
176                 xb, yb = data.get_batch()
177                 loss_np, _, summary = sess.run([loss, optim, merged_summary_op],
178                     feed_dict={x: xb, y: yb, keep_prob: drop_out})
179                 # logs every batch
180                 summary_writer.add_summary(summary, epoch * num_batches + i)
181                 avg_cost += loss_np/num_batches
182
183                 # Display logs per epoch step
184                 if (epoch+1) % display_epoch == 0:
185                     print("Epoch:", '%02d' % (epoch+1),
186                         "cost=", "{:.6f}".format(avg_cost))
187                 print('Validation Set Accuracy:',
188                     accuracy.eval({x: data.x_val, y: data.y_val, keep_prob: 1.0}))
189
190             # Test the model on separate data
191             print('Test Set Accuracy:',
192                 accuracy.eval({x: data.x_test, y: data.y_test, keep_prob: 1.0}))
193
194             print("Run the command line:\n--> tensorboard --logdir=./tf_logs ")
195
196     # ----- MODEL FUNCTIONS -----
197
198     # convolution layer: conv-> +bias -> activation -> pool
199     def conv_layer(x, W, b, cp):
200         x = tf.nn.conv2d(x, W, strides=[1, cp['ks'], cp['ks'], 1],
201             padding=cp['kpad'])
202         x = tf.add(x,b)
203         mean, var = tf.nn.moments(x, axes=[0,1,2])
204         x = tf.nn.batch_normalization(x, mean, var, 0, 1, .001)
205         x = tf.nn.relu6(x)
206         return tf.nn.avg_pool(x, ksize = [1, cp['p'], cp['p'], 1],
207             strides = [1, cp['ps'], cp['ps'], 1], padding=cp['ppad'])
208
209     # fully connected layer
210     def fc_layer(x, W, b):
211         x = tf.add( tf.matmul( tf.reshape(x, [-1, tf.shape(W)[0]]), W ), b)

```

```
212     mean, var = tf.nn.moments(x, axes=[0])
213     x = tf.nn.batch_normalization(x, mean, var, 0, 1, .001)
214     return tf.nn.relu6(x)
215
216 if __name__ == '__main__':
217     main()
```

Program: cifar100.py

```
1 import os
2 import pickle
3 import numpy as np
4 from tqdm import tqdm
5 import tensorflow as tf
6 import matplotlib.pyplot as plt
7 from tensorflow.keras.datasets.cifar100 import load_data
8
9 logs_path = "./tf_logs/"
10
11 input_size = 32 # 32x32 imgs
12 num_channels = 3 # RGB
13 num_classes = 100
14
15 # hyper-parameters
16 BATCH_SIZE = 200
17 NUM_EPOCHS = 15
18 display_epoch = 1
19 LEARNING_RATE = 5
20 # regularization parameters
21 drop_prob = 0.1
22 reg_scale = 1e-6
23
24 # https://stackoverflow.com/questions/38592324/one-hot-encoding-using-numpy/38592416
25 def get_one_hot(targets, nb_classes):
26     res = np.eye(nb_classes)[np.array(targets).reshape(-1)]
27     return res.reshape(list(targets.shape)+[nb_classes])
28
29 class Data(object):
30     def __init__(self):
31         np.random.seed(31415)
32         (self.x_train, y_train), (x_test, y_test) = load_data()
33
34         (self.x_train, x_test) = (self.x_train / 255.0, x_test / 255.0)
35
36         [(self.x_val, self.x_test), (y_val, y_test)] = \
37             [np.split(var,2) for var in [x_test, y_test]]
38
39         [self.y_train, self.y_val, self.y_test] = \
40             [np.squeeze(get_one_hot(y, num_classes)) for y in [y_train,y_val,y_test]]
41
42         self.index = np.arange(self.x_train.shape[0])
43
44     def get_batch(self):
45         choices = np.random.choice(self.index, size=BATCH_SIZE)
46         return self.x_train[choices,:,:, :], self.y_train[choices,:]
47
48 # ----- MODEL FUNCTIONS -----
49
50 # --- CONV LAYER WRAPPER --- w/ L2 regularization
51 # conv -> dropout -> BN -> relu -> max_pool
52 def conv_layer(input, filters, kernel_size, strides=2, is_training=True):
53     x = tf.layers.conv2d(
54         input, filters, kernel_size, strides=strides, padding='same',
55         kernel_regularizer=tf.contrib.layers.l2_regularizer(scale=reg_scale)
56     )
57     x = tf.layers.batch_normalization(x, training=phase, renorm=False)
58     x = tf.nn.relu6(x)
59     return x
60
61 # --- FULLY CONNECTED LAYER WRAPPER ---
62 # matmul -> dropout -> BN -> relu
63 def fc_layer(input, units, is_training=True):
64     x = tf.layers.dense(
65         input, units,
66         kernel_regularizer=tf.contrib.layers.l2_regularizer(scale=reg_scale)
67     )
68     x = tf.layers.batch_normalization(x, training=is_training, renorm=False)
69     x = tf.nn.relu6(x)
70     x = tf.layers.dropout(x, rate=drop_prob, training=is_training)
71     return x
```

```

72
73 x = tf.placeholder(tf.float32, [None, input_size, input_size, num_channels])
74 y = tf.placeholder(tf.float32, [None, num_classes])
75 phase = tf.placeholder(tf.bool) # is_training
76 lr = tf.placeholder(tf.float32) # is_training
77
78 # ---- NN ----
79
80 def f(x):
81     x = conv_layer(x, 32, 3, strides=2, is_training=phase)
82     print(x.get_shape())
83     x = conv_layer(x, 32, 3, strides=2, is_training=phase)
84     print(x.get_shape())
85
86     x = tf.layers.max_pooling2d(x, 3, 2, padding='same')
87     print(x.get_shape())
88
89     x = conv_layer(x, 32, 3, is_training=phase)
90     print(x.get_shape())
91     x = conv_layer(x, 32, 3, is_training=phase)
92     print(x.get_shape())
93
94     x = tf.layers.max_pooling2d(x, 3, 2, padding='same')
95     print(x.get_shape())
96
97     x = tf.layers.flatten(x)
98     x = fc_layer(x, 64, is_training=phase)
99     x = tf.layers.dense(x, num_classes)
100     return x
101
102 # models
103 logits = f(x)
104 prediction = tf.nn.softmax(logits)
105
106 correct_pred = tf.equal(tf.argmax(prediction, 1), tf.argmax(y, 1))
107 accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32))
108
109 # FROM KAROL
110 trash, guess_5 = tf.nn.top_k(prediction, k=5)
111 actual = tf.argmax(y, axis=1)
112 actual = tf.transpose(tf.stack([actual, actual, actual, actual, actual]))
113 correct_pred_5 = tf.equal(guess_5, tf.cast(actual, tf.int32))
114 accuracy_5 = 5.0*tf.reduce_mean(tf.cast(correct_pred_5, tf.float32))
115
116 # LOSS
117 update_ops = tf.get_collection(tf.GraphKeys.UPDATE_OPS)
118 with tf.control_dependencies(update_ops):
119     loss = tf.reduce_mean( tf.losses.softmax_cross_entropy(y, logits) ) \
120         + tf.reduce_mean( tf.losses.get_regularization_loss() )
121     optim = tf.train.GradientDescentOptimizer(learning_rate=lr).minimize(loss)
122
123 init = tf.global_variables_initializer()
124
125 # Create a summary to monitor cost tensor
126 tf.summary.scalar("loss", loss)
127 # Create a summary to monitor accuracy tensor
128 tf.summary.scalar("accuracy", accuracy)
129 # Merge all summaries into a single op
130 merged_summary_op = tf.summary.merge_all()
131
132 with tf.Session() as sess:
133     data = Data()
134     sess.run(init)
135     # op to write logs to Tensorboard
136     summary_writer = tf.summary.FileWriter(logs_path,
137         graph=tf.get_default_graph())
138     learning_rate = LEARNING_RATE
139     # training
140     for epoch in range(NUM_EPOCHS):
141         avg_cost = 0.
142         num_batches = int( np.ceil( data.index[-1] / BATCH_SIZE ) )
143
144         if epoch%3 == 0:

```

```

145         learning_rate = learning_rate/100
146
147     for i in tqdm(range(num_batches)):
148         xb, yb = data.get_batch()
149         loss_np, _, summary = sess.run([loss, optim, merged_summary_op],
150             feed_dict={x: xb, y: yb, phase: True, lr: learning_rate})
151         # logs every batch
152         summary_writer.add_summary(summary, epoch * num_batches + i)
153         avg_cost += loss_np/num_batches
154
155     # Display logs per epoch step
156     if (epoch+1) % display_epoch == 0:
157         print("Epoch:", '%02d' % (epoch+1),
158             "cost=", "{:.6f}".format(avg_cost))
159     print('TOP-1 Validation Set Accuracy:',
160         accuracy.eval({x: data.x_val, y: data.y_val, phase: False}))
161
162     print('TOP-5 Validation Set Accuracy:',
163         accuracy_5.eval({x: data.x_val, y: data.y_val, phase: False}))
164
165     # Test the model on separate data
166     print('TOP-1 Test Set Accuracy:',
167         accuracy.eval({x: data.x_test, y: data.y_test, phase: False}))
168
169     print('TOP-5 Test Set Accuracy:',
170         accuracy_5.eval({x: data.x_test, y: data.y_test, phase: False}))
171
172     print("Run the command line:\n--> tensorboard --logdir=./tf_logs ")

```