

Goal

This assignment was an attempt to replicate the results of the paper: Learning to See in the Dark [1], in which the group trains a fully convolutional neural network to infer a properly exposed RGB image from an input, severely underexposed, RAW image. The Seeing in the Dark (SID) group's code can be seen at <https://github.com/cchen156/Learning-to-See-in-the-Dark>

Remarks

RAW files are large. I decided to decimate my input RAW files and ground truth PNGs by a factor of 16 so that all the training images could be easily loaded into memory. It is possible that this decimation ruins the set up that makes the SID group's results possible, namely, the spatial context from the original RAW file may be lost.

Sad (approaching not sad) Results

After 100+ epochs, my network stagnated, producing only brown squares. I then changed some hyper parameters around, including input patch-size, activation function, kernel initialization, and annealed learning rates. I was then able to train to a much lower average cost, but the model still stagnates at a relatively high value. Further investigation is needed to train further. My next step, if time was permitting, would be to apply further data-augmentation (as performed in the paper). **Note:** input images are not shown as they are essentially black to the human eye.




Figure 1: 20 epochs, activation fcn: relu6, patch-size: 32




Figure 2: 40 epochs, activation fcn: relu, xavier init, patch-size: 64



Figure 3: 100 epochs, activation fcn: relu, xavier init, patch-size: 64

References

- [1] Chen Chen, Qifeng Chen, Jia Xu, and Vladlen Koltun, "Learning to See in the Dark", in CVPR, 2018. <https://arxiv.org/pdf/1805.01934.pdf>

Program: sid2.py

```
1  #!/usr/bin/python3
2  import time
3  import scipy as sp
4  import scipy.misc
5  import matplotlib.pyplot as plt
6  import pandas
7  import imageio
8  import numpy as np
9  import tensorflow as tf
10 from tqdm import tqdm
11 import rawpy
12 import os, glob
13 os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'
14
15 SIZE_REDUCTION = 16
16 logs_dir = './logs/'
17 tflog = logs_dir + 'tflog'
18 logfile = logs_dir + 'log.txt'
19 check_point_dir = './ckpt/'
20 in_dir = 'DWN%d/short/'%SIZE_REDUCTION
21 gt_dir = 'DWN%d/long/'%SIZE_REDUCTION
22
23 # hyper-parameters
24 BATCH_SIZE = 1
25 NUM_EPOCHS = 4000
26 SAVE_RATE = 20
27 DISP_RATE = 1
28 LR_DECAY_RATE = 40
29 LEARNING_RATE = 1e-3
30 PATCH_SIZE = 64
31
32 # -----
33 # ----- NETWORK -----
34 # -----
35
36 # activation fcn used in architecture
37 def act(x):
38     return tf.nn.relu(x)
39
40 def conv(x,chan,ksz):
41     return tf.layers.conv2d(x,chan,ksz,padding='SAME',activation=act,
42                             kernel_initializer=tf.contrib.layers.xavier_initializer())
43
44 # contraction module of UNET
45 def contract_mod(x,chans=None):
46     num_fchans = 2*int(x.shape[-1]) if chans==None else chans
47     c = conv(x,num_fchans,3)
48     c = conv(c,num_fchans,3)
49     p = tf.layers.max_pooling2d(c,2,2,padding='SAME')
50     return c,p
51
52 # upsample in1, concatenate both inputs
53 # output 2*c feature channels
54 def upsample_cat(in1,in2,c):
55     ups = tf.layers.conv2d_transpose(in1,c,2,2,padding='SAME')
56     x = tf.concat([ups,in2],3)
57     return x
58
59 # expansion module of UNET
60 # input1 upsample connection, input2 concat connection
61 def expand_mod(input1,input2):
62     out_chans= input2.get_shape()[-1]
63     x = upsample_cat(input1,input2,out_chans)
64     x = conv(x,out_chans,3)
```

```

65     x = conv(x,out_chans,3)
66     return x
67
68 # UNET ARCHITECTURE
69 def f(x):
70     # contract
71     c1,p1 = contract_mod(x,chans=32)
72     c2,p2 = contract_mod(p1)
73     c3,p3 = contract_mod(p2)
74     c4,p4 = contract_mod(p3)
75     c5,_ = contract_mod(p4)
76     # expand
77     e1 = expand_mod(c5,c4)
78     e2 = expand_mod(e1,c3)
79     e3 = expand_mod(e2,c2)
80     e4 = expand_mod(e3,c1)
81     # out of UNET
82     sub = tf.layers.conv2d(e4,12,1,padding='SAME')
83     # depth2space == subpixel reconstruction
84     # out chans = 12/(2*2) = 3 <-- perfect!
85     out = tf.depth_to_space(sub,2)
86     return out
87
88 # learning rate
89 lr = tf.placeholder(tf.float32)
90 # input, truth, est.
91 x = tf.placeholder(tf.float32, shape=[None,None,None,4])
92 y = tf.placeholder(tf.float32, shape=[None,None,None,3])
93 y_hat = f(x)
94
95 # -----
96 # ----- FILES -----
97 # -----
98
99 gt_fns = glob.glob(gt_dir+'*')
100 in_fns = glob.glob(in_dir+'*')
101 train_ids = [int(os.path.basename(fn)[0:5]) for fn in gt_fns]
102 [ x.sort() for x in [gt_fns, in_fns, train_ids] ]
103 # exposure ratios
104 ratios = np.empty((len(train_ids)))
105 for i in range(len(train_ids)):
106     ratios[i] = float(os.path.basename(gt_fns[i])[9:-5]) / \
107         float(os.path.basename(in_fns[i])[9:-5])
108 print(ratios[0])
109 print('Loading data...')
110 start_time = time.time()
111
112 gt_imgs = np.stack( [scipy.misc.imread(x) for x in gt_fns] )
113 in_imgs = np.stack( [np.load(x) for x in in_fns] )
114 # subtracting black level and normalizing to [0,1] scale
115 in_imgs = np.maximum(in_imgs-512,0) / (16383-512)
116
117 time_elapsed = time.time() - start_time
118 print('%3fs to load data'%time_elapsed)
119
120 # demo of batches
121 # yb = gt_imgs[0,:,:,:]
122 # xb = in_imgs[0,:,:,:]
123 # print(xb.shape,yb.shape)
124 # figure, (ax1,ax2) = plt.subplots(1,2)
125 # ax1.imshow(np.squeeze(xb))
126 # ax2.imshow(np.squeeze(yb))
127 # plt.show()
128
129 # -----
130 # ----- DATA -----

```

```

131 # -----
132
133 # pk: input image in pack form
134 # gt: ground truth image (in sRGB)
135 # ps: patch size (square)
136 # a patch of the packed bayers corresponds to a patch
137 # of the full size image (gt) that is twice as large
138 def to_patch(pk,gt,ps=PATCH_SIZE):
139     H = pk.shape[0]
140     W = pk.shape[1]
141     h = np.random.randint(H-ps)
142     w = np.random.randint(W-ps)
143     pk_patch = pk[h:(h+ps), w:(w+ps), :]
144     gt_patch = gt[2*h:2*(h+ps), 2*w:2*(w+ps),:]
145     return pk_patch, gt_patch
146
147 def batch_to_patch(xb,yb,ps=PATCH_SIZE):
148     xp = []
149     yp = []
150     H = xb.shape[1]
151     W = xb.shape[2]
152     for i in range(xb.shape[0]):
153         h = np.random.randint(H-ps)
154         w = np.random.randint(W-ps)
155         xp.append(xb[i, h:(h+ps), w:(w+ps), :])
156         yp.append(yb[i, 2*h:2*(h+ps), 2*w:2*(w+ps),:])
157     return np.stack(xp), np.stack(yp)
158
159 def get_batch():
160     choices = np.random.choice(len(train_ids),size=BATCH_SIZE)
161     rb = ratios[choices] # batch of ratios
162     xb, yb = batch_to_patch(in_imgs[choices,:,:,:], gt_imgs[choices,:,:,:])
163     for i in range(BATCH_SIZE):
164         xb[i,:,:,:] = xb[i,:,:,:]*rb[i]
165     return xb, yb
166
167 # demo of batches
168 # xb,yb = get_batch()
169 # print(xb.shape,yb.shape)
170 # figure, (ax1,ax2) = plt.subplots(1,2)
171 # ax1.imshow(np.squeeze(xb[0,:,:,:]))
172 # ax2.imshow(np.squeeze(yb[0,:,:,:]))
173 # plt.show()
174
175 # -----
176 # ----- TRAINING -----
177 # -----
178
179
180 # L1 LOSS ONLY
181 loss = tf.reduce_mean( tf.abs(y_hat - y) )
182 optim = tf.train.AdamOptimizer(learning_rate=lr).minimize(loss)
183 init = tf.global_variables_initializer()
184
185 # Create a summary to monitor cost tensor
186 tf.summary.scalar("loss", loss)
187 # Create a summary to monitor accuracy tensor
188 # tf.summary.scalar("accuracy", accuracy)
189 # Merge all summaries into a single op
190 merged_summary_op = tf.summary.merge_all()
191
192 with tf.Session() as sess:
193     sess.run(init)
194
195     # op to write logs to Tensorboard
196     summary_writer = tf.summary.FileWriter(tflog,

```

```

197         graph=tf.get_default_graph())
198     learning_rate = LEARNING_RATE
199     # training
200     st = time.time()
201     for epoch in range(NUM_EPOCHS):
202         avg_cost = 0.
203         num_batches = int( np.ceil( len(train_ids) / BATCH_SIZE ) )
204
205         if (epoch+1)%LR_DECAY_RATE == 0:
206             learning_rate = learning_rate*.1
207
208         for i in tqdm(range(num_batches)):
209             xb, yb = get_batch()
210             fd = {x: xb, y: yb, lr: learning_rate}
211             output, loss_np, _, summary \
212                 = sess.run([y_hat, loss, optim, merged_summary_op], feed_dict=fd)
213             # logs every batch
214             summary_writer.add_summary(summary, epoch * num_batches + i)
215             avg_cost += loss_np/num_batches
216
217             # save output of part of batch
218             if (epoch+1) % SAVE_RATE == 0:
219                 cat = np.concatenate((output[0,:,:,:],yb[0,:,:,:]),axis=1)
220                 scipy.misc.imsave(logs_dir + '%04d.jpg'%(epoch+1), cat)
221
222             # log to stdout at each epoch
223             if (epoch+1) % DISP_RATE == 0:
224                 print('Time: %06f, Epoch: %03d, Cost: %5ld'%(time.time()-st,epoch+1,
225                     avg_cost))
226
227             # print('Validation Set Accuracy:',
228                 # accuracy.eval({x: data.x_val, y: data.y_val, phase: False}))
229
230     print("Run the command line:\n--> tensorboard --logdir=./tf_logs ")

```